# Slide 1: Problem Introduction, Impact & User Breakdown

## THE GROUP ORDERING PAIN POINTS

- **The Endless Debate:** 73% of groups take 15+ minutes just to decide WHERE to order from because consensus is hard to reach.
- **Forgotten Preferences:** Dietary restrictions (e.g. allergies, pure veg) get overlooked, leading to wasted food and guest frustration.
- **The Budget Mismatch:** Different budgets create awkward ordering dynamics, with junior members often overpaying.
- **The Default Dictator:** One host takes over, resulting in individual choices that leave other members unsatisfied.

## BUSINESS & USER IMPACT METRICS

| Metric | Current State | Problem Impact |
|---|---|---|
| Group order abandonment | ~40% | Lower AOV & high checkout cart drops |
| Average decision time | 15-25 min | Ordering from safe, repetitive choices |
| Member satisfaction | ~55% | At least one member is unhappy with the choice |
| Repeat group orders | 2.3x/month | Suboptimal frequency due to high friction |

**Note: High abandonment directly reduces Swiggy's potential AOV conversion rate and increases platform acquisition costs.*

## CORE ASSUMPTIONS & HYPOTHESES

- **H1: Decision Speed:** AI-mediated group ordering will reduce final checkout decision times by 70%+.
- **H2: Trust & Satisfaction:** Transparent AI consensus will increase group satisfaction scores from 55% to 85%+.
- **H3: Commercial Value:** Smart group recommendation cards will drive a 25-35% uplift in Average Order Value (AOV).

*Insight: Shifting the cognitive load of decision-making from the user to the AI agent eliminates choice paralysis completely.*

## TARGET USER SEGMENTS BREAKDOWN

| Segment | TAM Share | Pain Level | Key Scenario & Need |
|---|---|---|---|
| Office Teams | 35% Share | **High** | Daily office lunch for 4-8 colleagues |
| Friend Groups | 40% Share | Medium-High | Weekend hangouts, parties & movie nights |
| Families | 20% Share | Medium | Dinners with strict multi-generation diet gaps |
| PG / Hostels | 5% Share | **High** | Budget-conscious students sharing bills |

*Key Focus: GenAI consensus caters specifically to the Office and Friend segments, which represent 75% of the total addressable market.*

# Slide 2: Why Generative AI?

## TECHNOLOGY COMPARISON: TRADITIONAL SYSTEMS VS. GENERATIVE AI

| Capability | Traditional Rules-Based Systems (Static ML) | Generative AI Agent Approach (Consensus Engine) |
|---|---|---|
| Preference Capture | Relies on fixed filters, tags, and static dropdown menus. Cannot parse unstructured custom requests or descriptive cravings. | Natural Language Processing parses complex sentences like 'I want something spicy but not too heavy' or 'comfort food' directly. |
| Dietary Constraints | Limits selections strictly to binary tags (e.g. Veg/Non-veg). Ignores complex ingredient exclusions. | Context-aware: Recognizes granular constraints like 'No onion-garlic, but dairy cheese and ghee are perfectly fine for us'. |
| Consensus Building | Imposes standard majority voting or default hosts' choices, creating unhappy members and higher cart drop rates. | Intelligent multi-objective optimization: weighs constraint matches across all users to select the most compatible venue. |
| Decision Transparency | No explanations provided. Users cannot see why a specific restaurant was suggested, leading to trust gaps. | Generates transparent, personalized natural explanations: 'Rahul gets Spicy Chicken Biryani, Priya gets Veg Biryani'. |

## WHY TRADITIONAL APPROACHES ARE INSUFFICIENT

- **Static Interfaces:** Users experience click-fatigue when forced to fill out repetitive multi-step forms just to express dinner cravings.
- **Collaborative Filtering Limits:** Traditional recommendation models only optimize for individual habits, completely failing in dynamic group dynamics.
- **Zero-compromise Voting:** Standard voting creates winners and losers, leading to cart dropouts or friction among friends.
- **Form Fatigue:** Static forms add friction and fail to capture conversational, contextual, or mood-based preference nuances.

## THE VALUE PROPOSITION OF GENERATIVE AI

- **Nuanced Semantic Search:** GenAI reads between the lines of human craving declarations, matching dish descriptions to complex user moods.
- **Conversational Interface:** Replaces complex surveys with simple chat loops, making input fast, social, and intuitive.
- **Soft & Hard Constraint Balancing:** Weighs dietary needs as hard limits while treating budget thresholds as soft guidelines.
- **Transparent Reasoning:** Surfacing AI explanations (why this venue is recommended) reduces group friction and builds strong system trust.

# Slide 3: User Segments & Solution Overview

## TARGET PERSONAS & DOCK PROFILES

- **Office Teams (35% Share):** Colleagues ordering daily. Host rotation is highly frustrating; selection is slow due to diverse dietary restrictions. Require speed, fairness, and variety.
- **Friend Groups (40% Share):** Friends ordering for weekend movie nights. Social friction from long WhatsApp debates and decision delays. Require an interactive, engaging flow.
- **Families (20% Share):** Multi-generational families ordering dinners. Severe preference splits between health-oriented parents and kids' fast-food cravings. Require safety checks.
- **PG & Hostel Groups (5% Share):** Roommates splitting weekend meals. Extremely budget-conscious. Require transparent per-person splitting to avoid payment friction.

*Strategy: By mapping conversational preferences to strict dietary tags, the AI Consensus engine meets the safety needs of families and the speed demands of office teams.*

## CORE DESIGN INSIGHT

### THE GROUP CONSENSUS PHILOSOPHY

*"The best group order isn't the one everyone loves - it is the one that no one hates. GenAI minimizes group friction rather than maximizing individual bias."*

## AI GADGETS TOOLKIT (DORAEMON METAPHOR)

- **Consensus Bread (AI Recommendations):** Processes all members' preferences to output the optimal dining venue matching everyone's constraints.
- **Cravings Megaphone (NLP Cravings Parser):** Translates unstructured speech (like 'something spicy but light') into clean catalog search tags.
- **Group Cart Copier (Auto-Cart Assembly):** Automatically populates individual items for each user, removing checkout errors.
- **Equal Split Wallet (Split-Bill System):** Dynamically splits final totals (taxes, fees, and items) to resolve reimbursement arguments.

# Slide 4: Solution Deep Dive & Product Flows

## STEP-BY-STEP USER INTERACTION FLOW

| 1. CREATE GROUP | --> | 2. ADD MEMBERS | --> | 3. AI COLLECT PREFS | --> | 4. AI CONSENSUS | --> | 5. CART / BILL SPLIT |
|---|---|---|---|---|---|---|---|---|
| [illegible] | | | | | | | | |

## THE AI INTERACTION LAYER

- **Semantic Cravings Parsing:** AI reads unstructured phrases (e.g. 'something hot and soupy') and maps them directly to menu item tags.
- **Hard Constraints Filtering:** Dietary boundaries (allergies, pure veg, religious rules) act as absolute binary search filters to keep recommendations 100% safe.
- **Conflict Resolution Algorithm:** If user preferences conflict (e.g. pizza vs curry), the AI searches menus containing top-rated items for both tastes.
- **Dynamic Persona Weights:** Weighs individual member requests against group profiles to prevent repetitive restaurant choices over time.

## KEY PRODUCT FEATURES & UX ENHANCEMENTS

- **Smart Quick Replies:** AI dynamically offers context-aware replies matching user order histories, allowing one-click preferences submission.
- **Satisfaction Compatibility Bars:** Clean UI indicators show users exactly how compatible each restaurant suggestion is with each member's taste.
- **Auto-Cart Construction:** AI automatically adds selected dishes for all group members, completely eliminating manual order coordination errors.
- **Transparent Split Receipt:** Displays individual breakdowns (subtotal, delivery, taxes) and split contributions directly on the final payment page.

# Slide 5: Success Metrics & Business Goals

## PRIMARY KEY PERFORMANCE INDICATORS (KPIS)

| Success Metric | Current Baseline | Target (3m) | Business Impact |
|---|---|---|---|
| Group Order AOV | Rs. 650 avg | Rs. 850 (+31%) | Higher basket sizes and revenue conversion |
| Decision Time | 18 min avg | 4 min avg (-78%) | Lower checkout abandonment rates |
| Order Frequency | 2.3x/month | 4.1x/month (+78%) | Increased repeat rates and customer LTV |
| Member Satisfaction | 55% satisfied | 88% satisfied | Boosts organic referral rates and virality |
| Funnel Abandonment | 40% drops | 12% drops (-70%) | Improves conversion rate metrics drastically |

**Note: baseline values are pulled from traditional Swiggy group cart checkouts without conversational consensus.*

## SECONDARY PRODUCT METRICS

- **AI Recommendation Acceptance:** Aim for >75% of groups to accept the top AI recommendation without manual overrides.
- **Invite Link Click Conversion:** Target >50% of shared invite link recipients to join the active group session.
- **Split Bill Payment Time:** Target all members completing equal split payments within 5 minutes of order placement.

## BUSINESS REVENUE GENERATION STRATEGY

- **Consolidated AOV Expansion:** Group carts aggregate multiple individual orders, driving higher average order values and delivery efficiencies.
- **Sponsored Consensus Ads:** Restaurants pay premium rates for transparent placement inside the AI recommendations grid.
- **Swiggy One Renewals:** Bundle AI Group Ordering into Swiggy One subscriptions to drive renewal rates and lower user churn.

# Slide 6: Pitfalls, Mitigations & Workarounds

## RISK MATURITY ASSESSMENT & MITIGATION STRATEGY

| Risk Concern | Risk Level | Mitigation Strategy | Technical Workaround |
|---|---|---|---|
| AI Hallucinations | High | Constrain recommendations strictly to active menus. | Implement schema validation checks that block non-existent items. |
| Privacy & Consent | Medium | Keep dietary/health profiles local and encrypted. | Enforce data encryption at rest; auto-wipe preference histories post-checkout. |
| Response Latency | Medium | Keep response delays below 2 seconds to prevent user dropouts. | Pre-fetch compatibility metrics; run consensus processing asynchronously. |
| User Trust Barriers | Medium | Surfaces the transparent logic behind each recommendation. | Displays a 'Why this?' breakdown mapping each member's needs explicitly. |
| Recommendation Bias | Medium | Audit recommendation scores to prevent brand favoritism. | Equally weight all compatible restaurants before applying sponsored overlays. |
| Adoption Friction | Low | Provide single-click guest ordering via web link. | Allow guests to join group sessions on web view without app downloads. |

## DEFENSE-IN-DEPTH FAILSAFE SYSTEM

To ensure 100% operational reliability, the GenAI integration uses a multi-layered safety net:

- **Layer 1 (Hard Boundaries):** The LLM operates on a highly constrained prompt scope, validating all outputs against live API menu items.
- **Layer 2 (User Overrides):** The host has final veto power. Any AI consensus card can be manually altered, overridden, or rebuilt.
- **Layer 3 (Feedback Loops):** Post-order satisfaction ratings are logged to continuously optimize recommendation weights.